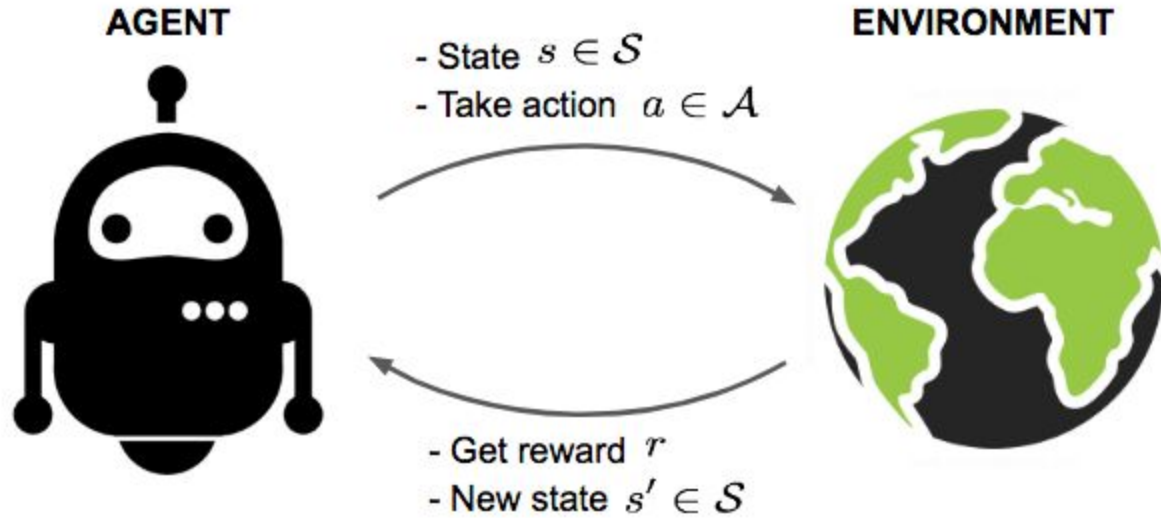


NeurIPS tutorial on RL and optimization

Part 1: RL as black-box optimization

Nicolas Le Roux, Google AI - Mila

Reinforcement learning



Agent is governed by a policy π

Image from Lilian Weng

The policy gradient objective

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_H)$$

$$G(\tau) = \sum_{h=0}^{H-1} \gamma^h r(s_h, a_h)$$

Q- and value functions

If I am in state s , take action a , then follow policy π , how much reward will I accumulate?

$$\begin{aligned} Q^\pi(s, a) &= r(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi(a'|s') r(s', a') + \dots \\ &= r(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi(a'|s') Q^\pi(s', a') \end{aligned}$$

$$V^\pi(s) = \sum_a \pi(a|s) Q^\pi(s, a)$$

Chapter 3.5 of [Reinforcement learning](#), Sutton and Barto

The policy gradient objective

$$V^{\pi}(s) = \sum_a \pi(a|s) Q^{\pi}(s, a)$$

$$V^{\pi}(\mu) = \mathbb{E}_{s_0 \sim \mu} [V^{\pi}(s_0)]$$

- μ Is the starting state distribution
- π Is the policy being optimized

Optimization in parameter space

In practice, π is parametrized by θ and we solve

$$\begin{aligned}\theta^* &= \arg \max_{\theta} V^{\pi_{\theta}}(\mu) \\ &= \arg \max_{\theta} J(\theta)\end{aligned}$$

This can be solved using standard gradient ascent

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\theta_t)$$

Optimization as an iterative procedure

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\theta_t)$$

$$\mathbb{E} [\| \nabla J(\theta) \|^2] \leq C_1 \frac{C_{\text{curv}}}{N} + C_2 \frac{\sigma}{\sqrt{N}}$$

Curvature is the limiting factor at first, then noise kicks in

Why does curvature hurt?

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\theta_t)$$

$$J(\theta) \approx J(\theta_t) + (\theta - \theta_t)^{\top} \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} \|\theta - \theta_t\|^2$$


- We make a small move because we do not trust the linear approximation
- How accurate is the approximation depends on how quickly the derivative changes

- If $\|\nabla_{\theta} J(\theta) - \nabla_{\theta} J(\theta')\| \leq L\|\theta - \theta'\|$


- Then

$$J(\theta) \geq J(\theta_t) + (\theta - \theta_t)^{\top} \nabla_{\theta} J(\theta_t) - \frac{L}{2} \|\theta - \theta_t\|^2$$

$$\theta_{t+1} = \theta_t + \frac{1}{L} \nabla_{\theta} J(\theta_t)$$



$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} \|\theta - \theta_t\|^2$$



Parameter linearity



Divergence

Newton method

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} \|\theta - \theta_t\|^2$$

Newton method

$$\begin{aligned}\theta_{t+1} &= \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} (\theta - \theta_t)^\top H(\theta - \theta_t) \\ &= \theta_t + \eta H^{-1} \nabla_{\theta} J(\theta_t)\end{aligned}$$

Newton method

$$\begin{aligned}\theta_{t+1} &= \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} (\theta - \theta_t)^\top H (\theta - \theta_t) \\ &= \theta_t + \eta H^{-1} \nabla_{\theta} J(\theta_t)\end{aligned}$$

Equivalent to standard gradient descent on $\theta' = H^{1/2}\theta$

Changing the divergence is equivalent to changing the representation

Mirror descent

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} \|\theta - \theta_t\|^2$$

Mirror descent

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} D_{\Phi}(\theta, \theta_t)$$

- D_{Φ} is called a *Bregman divergence*
- We want to find a divergence w/ good curvature properties

Reasoning with policies

- $$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} \|\theta - \theta_t\|^2$$

Reasoning with policies

- $$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

$$P_{\theta}(\tau) = \mu(s_0) \prod_{h=0}^H \pi_{\theta}(a_h | s_h) P(s_{h+1} | s_h, a_h)$$

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

- This problem does not have a closed form solution
- Two potential solutions:
 - Replace the KL with a quadratic approximation
 - Use a better approximation and do multiple optimization steps

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

- This problem does not have a closed form solution
- Two potential solutions:
 - Replace the KL with a quadratic approximation: NPG / TRPO
 - Use a better approximation and do multiple optimization steps: PPO

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

- Exactly minimizing the KL can be expensive
- Can we replace it with an approximation?

$$KL(P_{\theta_t} || P_{\theta}) = (\theta - \theta_t)^\top F(\theta_t)(\theta - \theta_t) + o(\|\theta - \theta_t\|^2)$$

- The minimization gives $\theta_{t+1} = \theta_t + \eta F(\theta_t)^{-1} \nabla_{\theta} J(\theta_t)$

Natural policy gradient (Kakade, 2001)

Replacing the penalty with a constraint

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

Replacing the penalty with a constraint

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) \text{ s.t. } KL(P_{\theta_t} || P_{\theta}) \leq C$$

- This does not immediately give a stepsize: line search!

$$\theta_{t+1} = \theta_t + \eta_t F(\theta_t)^{-1} \nabla_{\theta} J(\theta_t)$$

With η_t such that the constraint is satisfied

TRPO (Schulman et al., 2015)

From “ θ -linear” to “ π -linear”

$$\theta_{t+1} = \arg \max_{\theta} (\theta - \theta_t)^\top \nabla_{\theta} J(\theta_t) - \frac{1}{2\eta} KL(P_{\theta_t} || P_{\theta})$$

If the divergence deals with policies, why have an approximation linear in θ ?

From “ θ -linear” to “ π -linear”

$$J(\pi) = J(\pi_t) + \sum_{s,a} d^\pi(s) \pi(a|s) Q^{\pi_t}(s, a)$$

- $d^\pi(s)$ is the stationary distribution induced by π
- It has a complex derivative w.r.t. π

From “ θ -linear” to “ π -linear”

$$\begin{aligned} J(\pi) &= J(\pi_t) + \sum_{s,a} d^\pi(s) \pi(a|s) Q^{\pi_t}(s, a) \\ &\approx J(\pi_t) + \sum_{s,a} d^{\pi_t}(s) \pi(a|s) Q^{\pi_t}(s, a) \end{aligned}$$

Linear in π !

From “ θ -linear” to “ π -linear”

$$\pi_{t+1} = \arg \max_{\pi} \sum_{s,a} d^{\pi_t}(s) \pi(a|s) Q^{\pi_t}(s, a)$$

- Actions are sampled from $\pi_t(a|s)$
- We use the importance ratio $\frac{\pi_t(a|s)}{\pi(a|s)}$
- It is clipped to prevent poor estimation

PPO (Schulman et al., 2017)

The curse of deterministic policies

$$\mathbb{E}[\|\nabla J(\theta_N)\|^2] \leq C_1 \frac{C_{\text{curv}}}{N} + C_2 \frac{\sigma}{\sqrt{N}}$$

- Signal to noise ratio of deterministic policies is very small
- Second-order methods cannot address that
- We need to add curvature near the boundaries

Adding curvature near the boundaries

- $H(P_\theta) = - \sum_{\tau} P_\theta(\tau) \log P_\theta(\tau)$: entropy regularization, small effect
- $KL(P_{\theta_t} || P_\theta)$: local KL, slows down movement
- $KL(P_{\theta_0} || P_\theta)$: relative entropy, acts as a log barrier, faster rates

What about variance?

- Getting the true gradient is difficult
- Stochastic estimates of the gradient

$$\mathbb{E} \left[\|\nabla J(\theta)\|^2 \right] \leq C_1 \frac{C_{\text{curv}}}{N} + C_2 \frac{\sigma}{\sqrt{N}}$$

Control variates

- We want to estimate $\mathbb{E}_{\xi}[\nabla_{\theta}J(\theta, \xi)]$ from samples $\nabla_{\theta}J(\theta, \xi_i)$
- Imagine we know $\mathbb{E}_{\xi}[z(\xi)]$ for some variable z
- Then $\nabla_{\theta}J(\theta, \xi_i) - z(\xi_i) + \mathbb{E}_{\xi}[z(\xi)]$ has the correct expectation
- It also has lower variance if $z(\xi_i)$ is positively correlated with $\nabla_{\theta}J(\theta, \xi_i)$

Example: stochastic variance reduction methods

$$\theta_{t+1} = \theta_t - \alpha \left(\nabla_{\theta} J(\theta, \xi_i) - g_i + \frac{1}{N} \sum_j g_j \right)$$

- g_i is the gradient for datapoint i stored in memory
- We store the last computed gradient for each datapoint in memory
- This has (roughly) the same convergence rate as a batch method

Control variates in RL: baselines

$$\nabla_{\theta} J(\theta) = \sum_{s,a} d^{\pi}(s) \pi(a|s) Q^{\pi}(s,a) \nabla_{\theta} \log \pi(a|s)$$

$$\nabla_{\theta} J(\theta) = \sum_{s,a} d^{\pi}(s) \pi(a|s) \left(Q^{\pi}(s,a) - z(s) \right) \nabla_{\theta} \log \pi(a|s)$$

- $z(s)$ is called a (state-dependent) baseline
- A carefully chosen baseline can reduce the variance
- Other techniques exist but they bias the gradient (more w/ Martha)

Conclusion

- Policy gradient methods can be tackled as a pure optimization problem
- One needs to deal with curvature and noise
- Curvature: careful choice of the parametrization can help
- Entropy and relative entropy also add curvature
- Noise: baselines are an easy tool but they have limited impact